
MIBE: Multi-subject Interaction Benchmark and Evaluator for Personalized Image Generation

Anonymous Author(s)

Affiliation

Address

email

Abstract

Multi-subject personalized image generation requires the precise rendering of all requested reference identities and their specified interactions based on a guiding prompt. However, state-of-the-art models still struggle with this process, frequently omitting subjects, failing to preserve reference appearances, or misattributing interactions. Furthermore, existing metrics designed primarily for single-subject fidelity cannot reliably capture these errors, suffering severe degradation in ranking separability and failing to align with human preference as the subject count increases. To address this gap, we introduce Multi-subject Interaction Benchmark and Evaluator (MIBE), a unified framework comprising a Multi-subject Interaction Benchmark (MIB) and a Multi-subject Interaction Evaluator (MIE). MIB systematically covers diverse relation types and scene complexities through a decoupled data regime. This consists of a 60K-pair VLM-labeled Silver Set for scalable metric training and a 4K-pair double-blind Human Evaluation Gold Set covering a diverse range of state-of-the-art generators, with the Silver Set reaching 95.1% cross-VLM preference agreement. To demonstrate the utility of this benchmark, we present MIE, a lightweight, reference-conditioned evaluator trained exclusively on the Silver Set with a dual-head ranking and diagnosis objective. MIE exhibits strong cross-generator generalization on the Gold Set, achieving 0.922 overall pairwise accuracy against human preference, including 0.982 on seen generators and 0.884 on unseen generators. By outperforming a broad spectrum of baseline metrics, including CLIP and DINO variants, MIE demonstrates that diagnostic supervision can preserve ranking separability and human alignment where traditional evaluators collapse. Ultimately, by providing data with the diagnostic and preference labels required to train robust evaluators, MIBE serves as a comprehensive framework to benchmark multi-subject image generation and guide future model alignment.

1 Introduction

State-of-the-art personalized image generators can faithfully render a single reference identity, but as soon as a second subject is introduced, let alone four or eight, the generated image often contains the right scene, the right colors, and the wrong subjects. A model conditioned on reference photos of a specific child and a specific dog, asked to depict the child playing with the dog, may instead produce a generic child with a different dog, two children and no dog, or the right child standing beside an empty patch of grass. These are not merely aesthetic flaws; they are *binding* failures, and they are largely invisible to current metrics.

We call this the *binding* problem: every requested reference subject must (1) appear in the generated image (*Existence*), (2) preserve its reference appearance without leakage from other subjects (*Appearance*), and (3) participate in the prompt-specified role, relation, or object assignment (*Interaction*). These dimensions are coupled: a missing subject often causes the model to redistribute its

visual features onto surviving subjects, making appearance corruption a downstream consequence of existence failure (Appendix A). The problem is especially fragile in contact-rich and occlusion-heavy scenes, where local evidence is ambiguous. Prior stress-test evaluation has documented this “illusion of scalability,” showing identity collapse in dense scenes while global CLIP-based scores remain misleadingly stable [3]; MIBE complements that curated identity-collapse stress test with human-grounded labels, factorized binding stress, and a learned evaluator.

Existing automatic metrics are fragmented. CLIP- and DINO-based reference similarity captures appearance overlap but cannot verify which subject participates in which interaction. General preference models such as PickScore [7], ImageReward [24], and HPS [22] can favor visually polished but compositionally broken images. As Figure 1 shows, as subject count grows from 2 to 8, the pairwise agreement of standard metrics with double-blind human judgment drops toward random choice; on MIB-Gold, HPS v2.1 reaches only 0.520 agreement and PickScore falls slightly below random at 0.486. Existing benchmarks are similarly under-equipped, focusing on single-subject fidelity [15, 13], text-only composition [6], or broad multi-subject capability [2, 17], without factorizing the conditions that induce binding failures. Binding is therefore both a generation and a measurement bottleneck.




Figure 1: **Existing metrics approach random agreement as subject count grows; MIE retains higher human alignment.** Pairwise agreement with double-blind human preference on MIB-Gold by subject count ($N \in \{2, 4, 6, 8\}$). Standard metrics, including SCR [3], approach random agreement (dashed line) in high-subject-count settings; MIE, trained exclusively on MIB-Silver, retains higher agreement, with the gap widening as scenes grow more crowded.




Figure 2: Overview of MIBE. MIB constructs a controlled benchmark through reference pooling, hierarchical prompt construction, Silver Set supervision, and Gold Set human evaluation. MIE is trained on Silver labels and evaluated against double-blind human Gold labels for cross-generator, human-aligned assessment.

We address this gap with **MIBE**, a unified framework that pairs a controlled **Multi-subject Interaction Benchmark (MIB)** with a learned **Multi-subject Interaction Evaluator (MIE)**. As summarized in Figure 2, MIBE decouples benchmark construction from evaluator learning: MIB defines controlled prompt-reference tasks and silver/gold supervision, while MIE uses this supervision to learn human-aligned ranking and diagnostic evaluation.

In summary, this paper makes three contributions:

- 1) We introduce **MIB**, a controlled benchmark for reference-conditioned multi-subject generation that factorizes subject count, human/object composition, and relation type, with a 60K dual-VLM-consensus silver set and a 4,020-pair double-blind human-labeled gold set covering six state-of-the-art generators.
- 2) We propose **MIE**, a reference-conditioned evaluator that jointly learns pairwise ranking and structured diagnostics over *Existence*, *Appearance*, and *Interaction*, providing both human-aligned preference prediction and interpretable failure attribution.
- 3) Through cross-generator meta-evaluation on MIB-Gold, we reveal systematic failure patterns, failure rates that scale with subject count, coupled Existence–Appearance failures, and interaction-induced subject deformation, providing actionable targets for future model alignment.

2 Related Work

Personalized and Multi-Subject Generation Personalized image generation has evolved from early optimization-based methods like DreamBooth [15] and Textual Inversion [4] to efficient encoder-based models such as IP-Adapter [26] and PhotoMaker [10]. While effective for single subjects, these methods struggle with binding in multi-subject scenarios, ensuring each reference identity is correctly assigned to its specific role and attributes. Recent works like FastComposer [23], MS-Diffusion [18], and Interact-Custom [25] address this through localized attention and identity decoupling. However, a significant measurement gap remains: existing methods are rarely tested in contact-rich or occlusion-heavy regimes where subjects overlap, leading to frequent failures in physical grounding and identity leakage that current metrics fail to capture.

Benchmark Evolution While compositional benchmarks like T2I-CompBench [6] evaluate text-to-image alignment, they lack the reference conditioning necessary to measure identity preservation. Recent efforts such as XVerseBench [2], PSRBench [17], and MultiHuman-Testbench [1] have introduced multi-subject evaluations, yet they primarily focus on subject count or spatial positioning. Unlike MultiBind [16], which explores attribute-level misbinding, MIB factorizes complex physical interactions and occlusions. It serves as both a diagnostic testbed and a training substrate, decoupling scalable "silver" supervision from human "gold" evaluation to provide a more granular view of how models handle locally ambiguous visual evidence. A detailed comparison with prior benchmarks is provided in Appendix B.

Metrics and Scalable Evaluators Standard metrics often provide incomplete assessments of multi-subject quality: CLIP-based scores lack instance-specific grounding, while reward models like ImageReward [24] often prioritize aesthetics over logical binding. Although Vision-Language Models (VLMs) are increasingly used as automated judges, they remain sensitive to prompting and poorly calibrated for absolute scoring [8]. MIBE addresses these limitations by adopting a pairwise ranking format and a latent scalar utility model, rather than absolute numeric ratings. By combining an "errors-first" VLM consensus for scalable MIB-Silver labels with human-verified MIB-Gold data, MIBE unifies preference learning with structured diagnostic supervision for more robust, reference-grounded evaluation.

3 Methods

3.1 MIB Construction

MIB is the data substrate of MIBE, consisting of controlled prompt-reference tasks, paired generated candidates, and preference/diagnostic labels. Each task specifies a guiding prompt, a set of reference subjects, and two candidate generations produced under the same prompt-reference condition. This decoupled design separates reference construction, prompt-task design, candidate generation, and




Figure 3: Hierarchical prompt construction and entity-composition templates. Each Level-8 seed is reduced into strict Level-6/4/2 subsets by removing entities together with their associated relations, preserving core scene semantics while varying entity-relation density. The example illustrates a human-heavy physical-interaction bucket, while the table summarizes the human/non-human composition templates used across balanced, human-heavy, and object-heavy ratio types.

label acquisition, allowing subject identity, scene complexity, relation type, generator source, and annotation source to be controlled independently.

3.1.1 Prompt Construction

We construct 60K prompt-reference tasks from 15K hierarchical Level-8 seeds. Each seed starts from an eight-subject scene and is reduced into Level-6, Level-4, and Level-2 variants by removing entities together with their associated relations, as illustrated in Figure 3. This strict subset construction keeps scene semantics approximately invariant across levels, so performance degradation can be more directly attributed to increasing entity-relation density.

The prompt space is factorized into 36 buckets from 4 subject counts, 3 human/non-human ratio types, and 3 relation types. The entity-composition templates for balanced, human-heavy, and object-heavy ratios are shown in Figure 3. Relation types include non-contact relations, occlusion relations, and physical interactions. When relations overlap, we assign the task to the most restrictive category using the priority physical interaction > occlusion > non-contact. To keep role assignment interpretable, each entity participates in at most one core relation.

All prompts use a white-studio setting and a global realistic-scale constraint to reduce background, lighting, layout, and scale confounders. We apply rule-based parsing and a lightweight LLM validator to verify: (i) assigned entities are explicitly mentioned, (ii) unassigned entities are excluded, (iii) human/non-human counts match the target bucket, (iv) lower-complexity prompts remain strict subsets of higher-complexity prompts, (v) duplicate prompts are removed, (vi) no entity is reused as the subject of multiple independent relations, and (vii) interaction-specific keywords are excluded from non-contact prompts.

3.1.2 Reference and Candidate Image Generation

MIB uses a fixed pool of 80 GPT-Image-1 reference subjects, including 30 human and 50 non-human identities. Human references vary in age, gender presentation, ethnicity, clothing, hairstyle, and appearance, while non-human references include animals, furniture, tools, wearables, and manipulable objects. All references are standardized as isolated single-subject studio images at 512×512 , allowing the same identity anchors to be recombined across prompts. For each prompt-reference task, a personalized generator receives the guiding prompt and reference images as visual conditioning inputs, producing a candidate image. MIB forms pairwise comparisons by placing two candidates under the same prompt-reference condition, so labels capture relative generation quality rather than prompt difficulty. References are generated separately from all candidate generators to reduce generator-specific leakage.

For the Silver Set, we generate 60K Nano Banana-versus-MOSAIC candidate pairs. This matchup was selected to maximize failure diversity observed in preliminary inspection, including subject omission, identity loss, appearance leakage, and relation hallucination. These complementary errors yield informative pairwise supervision across Existence, Appearance, and Interaction. Reference images are generated in a separate single-subject construction stage, while candidate images are generated as multi-subject interaction scenes under prompt-reference conditioning. For MIB-Silver, the reference generator is also model-disjoint from the two candidate generators, reducing generator-specific leakage in scalable supervision.

MIB-Gold contains 4,020 human-labeled pairs: 2,520 pairs from two fixed cross-model matchups among GPT-Image-1.5, Flux2, Seedream 4.5, and GLM, plus 1,500 Nano Banana-versus-MOSAIC pairs for seen-generator evaluation. Candidate order is randomized, and invalid generations are removed before annotation. Dataset statistics are summarized in Table 1.

Attribute	Silver Set	Gold Set
Purpose	MIE training	Human benchmark evaluation
Scale	60,000 candidate pairs; 56,909 usable pairs	4,020 human-labeled pairs
Prompt coverage	Hierarchical prompts with strict high-to-low complexity subsets, approximately uniformly distributed across relation type, human/object ratio, and subject count.	
Candidate generators	Nano Banana, MOSAIC	GPT-Image-1.5, Flux2, Seedream 4.5, GLM; plus Nano Banana/MOSAIC for seen-generator evaluation
Annotation source	SOP-guided dual VLM consensus	Double-blind human annotation
Annotators	Gemini-2.5-Flash and Gemini-3.1-Flash-Lite	Two independent human annotators per pair
Label types	Pairwise preference; Existence; Appearance; Interaction	Pairwise preference; Existence; Appearance; Interaction
Evaluation role	Train and Val	Seen-generator: 1,500 pairs; unseen-generator: 2,520 pairs

Table 1: Dataset statistics for MIB. The Silver Set provides scalable SOP-guided VLM supervision for training MIE, while the Gold Set is fully human-labeled under a double-blind protocol.

3.1.3 Label Construction

MIB provides two complementary label types: pairwise preference labels for global ranking and diagnostic labels for failure localization. Pairwise preference asks which candidate is better overall under the same prompt-reference condition, avoiding the calibration difficulty of absolute scores while preserving relative severity information. Diagnostic labels decompose multi-subject generation quality into three dimensions:

- **Existence:** whether all requested reference subjects are present without omission, duplication, or severe unidentifiable collapse.
- **Appearance:** whether generated subjects preserve the structure and local appearance of their references without severe distortion or cross-subject feature bleeding.
- **Interaction:** whether subject-to-subject relations, actions, and object assignments align with the prompt.

For diagnostic labels, annotator disagreements are averaged as soft scores, while pairwise preferences require strict consensus.

MIB-Silver annotation with SOP-guided VLM judges. The SOP decomposes each comparison into Existence, Appearance, and Interaction checks and requires candidate-specific flaw logs before the final preference decision, grounding judgments in concrete failures rather than holistic aesthetics. We independently query `gemini-2.5-flash` and `gemini-3.1-flash-lite-preview` with the same SOP, retain pairs with matching pairwise preferences, and average their per-dimension diagnostic

167 outputs as soft supervision. Across 59,852 valid overlapping Silver comparisons, the judges reach
168 95.1% preference agreement, yielding 56,909 usable Silver pairs after agreement-based filtering. The
169 full SOP appears in Appendix C.

170 **VLM judge selection.** This dual-judge configuration was selected via a 100-pair pilot human-
171 alignment study: each judge alone reaches 98.9% pair-level alignment with human winners, and their
172 strict-consensus combination retains 86.8% of pilot pairs at human alignment.

173 **MIB-Gold annotation with double-blind human evaluation.** Each Gold pair is independently
174 reviewed by two annotators under a double-blind protocol, providing both pairwise preference and
175 Existence/Appearance/Interaction labels. After preference-consistency filtering, both Gold subsets
176 retain over 90% of pairs, indicating stable human judgments while preserving enough difficulty to
177 expose non-trivial generator and evaluator disagreement.

178 3.2 Evaluator Modeling

179 The primary goal of MIBE is to establish a rigorous evaluation framework for multi-subject personal-
180 ized image generation. To demonstrate the reusability of our benchmark and provide an immediate,
181 out-of-the-box tool for the community, we instantiate the **Multi-subject Interaction Evaluator**
182 **(MIE)**. Trained exclusively on MIB-Silver and assessed on MIB-Gold, MIE bridges the gap between
183 human judgment and automated metrics.

184 **Input-Output Formulation.** For each task, let p denote the prompt, $R = \{r_1, \dots, r_N\}$ denote the
185 set of reference images, and let y denote a candidate generation. MIE is a reference-conditioned
186 multimodal evaluator that maps the triplet (R, p, y) to two types of outputs:

- 187 1. A scalar quality score $s_\theta(R, p, y) \in \mathbb{R}$ used for global pairwise ranking.
- 188 2. A three-dimensional diagnostic prediction vector of logits $\hat{d}_\theta(R, p, y) = (\hat{d}_{\text{exist}}, \hat{d}_{\text{app}}, \hat{d}_{\text{inter}})$,
189 corresponding to *Existence*, *Appearance*, and *Interaction*.

190 **Dual-Head Design.** A purely diagnostic model with binary labels cannot capture relative failure
191 severity (e.g., missing one subject vs. missing all). Conversely, a purely scalar reward model acts as
192 an opaque black box, offering no actionable feedback on the nature of the failure. MIE solves this
193 via a dual-head architecture. The **ranking head** provides a continuous global signal for pairwise
194 comparison, which is essential for leaderboard benchmarking. Simultaneously, the **diagnostic head**
195 explicitly attributes errors to Existence, Appearance, or Interaction, providing interpretable feedback
196 for model developers. By optimizing these jointly ($\mathcal{L} = \alpha\mathcal{L}_{\text{rank}} + \beta\mathcal{L}_{\text{diag}}$), the scalar score is forced
197 to ground itself in concrete binding failures rather than superficial aesthetics. This design mirrors the
198 human annotation procedure directly, preserving both ranking separability and interpretability.

199 **Lightweight Tuning.** To ensure that the evaluator remains practical and highly reusable, we adopt
200 a parameter-efficient fine-tuning strategy (LoRA and layer-only updates) over off-the-shelf vision-
201 language backbones (e.g., Qwen3.5-VLM). This design choice proves the exceptional quality and
202 data efficiency of our dataset: one does not need to pretrain a massive multimodal model from scratch;
203 simply fine-tuning a lightweight VLM on just a 9.6K subset of the 56.9K available MIB-Silver pairs
204 yields an evaluator that highly correlates with humans. This guarantees that MIB-Silver is a highly
205 reusable resource for future metric alignment in the community.

206 4 Results

207 All training and evaluation are conducted on a single NVIDIA A100 GPU, with 16 vCPUs (Intel(R)
208 Xeon(R) Platinum 8462Y+) and 251 GB of system memory. The configurations of training and
209 evaluation processes are provided in Appendix E.

210 4.1 Existing Metrics Fail on MIB-Gold

211 We begin by testing whether existing automatic metrics can serve as reliable proxies for human
212 judgment on MIB-Gold. This is a meaningful stress test rather than an artifact of noisy annotation:
213 after preference-consistency filtering, MIB-Gold retains 94.1% of seen-generator pairs and 90.4%
214 of unseen-generator pairs, providing a stable yet challenging benchmark for metric evaluation. To

establish a comprehensive performance ceiling, we therefore select a representative suite of baselines spanning four distinct paradigms: 1) **Low-level Reconstruction Metrics**: PSNR [5], SSIM [19], and MS-SSIM [20] are included to measure pixel-level fidelity; 2) **Semantic Alignment Metrics**: DINOv2 [11], CLIP [14], and SigLIP [27] (both Text-based and Image-based) are used to evaluate high-level semantic consistency and identity similarity; 3) **General Human Preference Models**: PickScore [7], ImageReward [24], and HPS v2.1 [22], which are specifically trained on human feedback for text-to-image alignment and aesthetics; 4) **Identity-Specific Metrics**: SCR (Subject Consistency Rate) [3], which focus on identity-preserving features critical for personalization.

To assess whether existing metrics can reliably evaluate multi-subject personalized generation, we conduct a systematic comparison against human pairwise preference on MIB-Gold. Specifically, we evaluate each metric on 4,020 valid pairwise comparisons, measuring pairwise agreement with double-blind human annotations, where 0.5 corresponds to random chance in a balanced binary setting. As shown in Figure 1, the results reveal a consistent failure across existing paradigms. General-purpose preference models collapse near or below random chance, PickScore achieves 0.4857, HPS v2.1 reaches 0.5201, and PSNR yields 0.3987, indicating active misalignment with human judgment. Identity-specific metrics including SCR, DINOv2, and SigLIP-I demonstrate non-trivial alignment, yet their gains remain confined to single-dimensional facets such as identity preservation or global semantic consistency. Critically, no existing metric provides a unified reasoning signal across Existence, Appearance, and Interaction simultaneously—the multi-dimensional structure underlying human preference in this setting. This confirms a fundamental evaluation gap and motivates the design of MIE.

4.2 MIE Aligns Better with Human Preference

Given that existing metrics are misaligned with human preference on MIB-Gold, we next ask whether MIB-Silver can be used to train a better evaluator. Our results show that it can. Across all six exported checkpoints, the strongest variant is `qwen35_4b_lora_layer`, which achieves an overall pairwise accuracy of 0.922, including 0.982 on the seen-generator subset (Nano Banana/MOSAIC) and 0.884 on the unseen-generator subset (GPT-Image-1.5, Flux2, Seedream 4.5, and GLM). This substantially exceeds the strongest third-party baseline, showing that supervision derived from MIB can be translated into a materially stronger human-aligned metric.

The gains are not limited to pairwise ranking. The same 4B LoRA-layer model reaches a macro-F1 of 0.818, indicating that the evaluator is not merely learning a shallow preference score, but capturing a meaningful portion of the fine-grained diagnostic structure underlying human judgments. More importantly, this result shows not only that a larger model can rank better, but that MIB-Silver provides sufficiently structured supervision to train a human-aligned and diagnostically meaningful evaluator. As Figure 4 shows, the strongest model is also strongest on category-level F1, linking pairwise alignment to meaningful per-dimension supervision. This matters because the binding problem in multi-subject personalization is inherently multi-dimensional: a useful evaluator must jointly reason about existence, appearance, and interaction rather than collapse everything into a single weak aesthetic signal.

More broadly, the six-checkpoint comparison now supports a full scaling narrative. LoRA-layer variants consistently outperform their layer-only counterparts, and the largest LoRA-based model achieves the strongest overall human alignment. As shown in Figure 4, the advantage is visible not only in pairwise accuracy but also in the category-level F1 breakdown. Thus, MIB does not only expose the failure of existing metrics; it also enables the training of a new evaluator that tracks human preference far more faithfully while preserving meaningful per-dimension diagnostics in multi-subject settings.

MIE Breakdown Analysis We analyze where the gains of the learned evaluator actually come from. A consistent pattern emerges across all six checkpoints: every model performs better on the seen-generator subset than on the unseen-generator subset, confirming that cross-generator generalization remains a non-trivial challenge. However, the size of this gap depends strongly on the tuning regime. The smallest seen-to-unseen drop is achieved by `qwen35_4b_lora_layer` at -0.098 , whereas the largest drop appears in `2b_layer_only` at -0.221 . This shows that additional model capacity alone is not enough; how that capacity is adapted is equally important.




Figure 4: Human alignment of MIE variants on MIB-GoId. Left: overall, seen-generator, and unseen-generator pairwise accuracy. Right: category-level F1 across existence, appearance, and interaction. The 4B LoRA-layer evaluator is the strongest overall model and remains clearly above third-party baselines even on the unseen-generator subset.

The LoRA-versus-layer comparison reinforces this conclusion. At 2B, LoRA-layer tuning improves pairwise accuracy by 0.061 and macro-F1 by 0.116 relative to 2B `layer_only`. At 4B, it still yields gains of 0.046 in pairwise accuracy and 0.044 in macro-F1 over 4B `layer_only`. Even at 0.8B, where the pairwise gain is nearly neutral, LoRA-layer tuning improves macro-F1 by 0.128. These results suggest that the primary benefit of LoRA-layer tuning is not merely stronger ranking, but consistently sharper diagnostic discrimination. In other words, the full six-checkpoint comparison supports a benchmark-utility claim rather than a simple “bigger is better” story.

The category-level breakdown points to the same conclusion. Existence remains the easiest dimension overall, whereas Appearance and Interaction provide the more discriminative diagnostic challenge. As shown in Figure 5, generator shift is not uniform across dimensions: the strongest checkpoint remains highly stable on Existence, while the more semantic categories continue to differentiate model quality and adaptation strategy. Overall, the breakdown analysis clarifies that the best-performing evaluator is not simply the largest model, but the model that combines additional capacity with the right adaptation strategy and preserves diagnostic sensitivity under distribution shift.

5 Limitations

Scale and Diversity of the Gold Set. While our 4K gold evaluation set was curated under a rigorous double-blind protocol, it inherently captures a finite subspace of the vast combinations of subjects and interactions. The current benchmark relies on 80 unique reference subjects, which, despite covering major categories of humans, animals, and everyday objects, may not fully represent the long-tail distribution of real-world identities or culturally specific concepts. The white-studio setting and global realistic-scale constraint further narrow the visual distribution to controlled compositional stress, leaving in-the-wild backgrounds and stylistic variations beyond the current scope. Future iterations will expand the reference concept library along long-tail axes and extend evaluation to less constrained visual settings.

Temporal Snapshot of Baseline Generators The MIB gold set reflects the performance of six representative generators at the time of collection. Given the rapid velocity of the field, newer architectures may encounter unique failure modes, such as specific multi-modal reasoning bottlenecks, not fully captured by the current distribution. To ensure the benchmark’s longevity, we commit



Figure 5: Breakdown analysis of MIE variants. Left: seen-to-unseen generalization gap. Middle: LoRA-layer gains over layer-only tuning at different model scales. Right: category-level F1 for the strongest checkpoint. The results show that LoRA-layer tuning improves diagnostic quality across scales, while generator shift affects diagnostic categories non-uniformly: Existence remains the most stable dimension, and Appearance/Interaction remain more discriminative.

297 to releasing versioned extensions as transformative models emerge, providing a standardized re-
 298 annotation protocol to facilitate community-driven growth.

299 **VLM-Induced Noise in the Silver Set** The 60K silver set utilizes a VLM-based pipeline for
 300 scalable preference pseudo-labeling. Despite its utility, VLM labeling is susceptible to model-specific
 301 biases, such as the tendency to conflate surface-level image fidelity with structural binding accuracy.
 302 Furthermore, these models may underperform on culturally specific concepts underrepresented in
 303 their pretraining data. We recommend that practitioners using the silver set for model alignment
 304 (e.g., via DPO) treat these labels as high-fidelity but noisy proxies, incorporating human-in-the-loop
 305 auditing for safety-critical applications.

306 6 Conclusion

307 We introduced **MIBE**, a benchmark-and-metric framework for multi-subject personalized image gen-
 308 eration, centered on the concept of *binding*: the correct assignment of visual identities, appearances,
 309 and interactions to their designated subjects. The primary contribution of this work is **MIB**, a large-
 310 scale, rigorously annotated benchmark constructed via hierarchical prompt design and factorized
 311 scene buckets. MIB provides (i) a 60K-pair silver training corpus with VLM-derived preference
 312 labels and attribute annotations, reaching 95.1% cross-VLM preference agreement, enabling scalable
 313 metric learning and model alignment; and (ii) a 4K-pair gold evaluation set with double-blind human
 314 judgments across six state-of-the-art generators, providing a reproducible, human-grounded testbed
 315 for future comparisons. To our knowledge, MIB is the first dataset to diagnose binding failures
 316 across factorized relation types and subject counts up to eight. Building on MIB, we further propose
 317 **MIE**, a reference-conditioned evaluator that operationalizes binding quality into three interpretable
 318 dimensions, Existence, Appearance, and Interaction, alongside a scalar preference signal. MIE
 319 validates the utility of MIB as a training resource: trained exclusively on the silver set, MIE achieves
 320 0.922 overall pairwise accuracy against human preference on the gold set, including 0.982 on seen
 321 generators and 0.884 on unseen generators, without any target-specific fine-tuning, demonstrating that
 322 the annotation protocol captures transferable binding signals rather than generator-specific artifacts.
 323 We position MIBE as open infrastructure for the multi-subject personalization community, and hope
 324 it becomes a standard testbed that accelerates progress on the binding problem and raises the bar for
 325 rigorous evaluation in personalized image generation. All data, annotations, and evaluation code will
 326 be publicly released under permissive licenses to maximize reproducibility and community uptake.

References

- [1] Shubhankar Borse, Sungrae Choi, Sanghyun Park, Jihoon Kim, Sayak Kadambi, Risheek Garrepalli, Suyeon Yun, Munawar Hayat, and Fatih Porikli. Multihuman-testbench: Benchmarking image generation for multiple humans. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- [2] Hao Chen et al. Xverse: Consistent multi-subject control of identity and semantic attributes via dit modulation. *arXiv preprint arXiv:2506.21416*, 2025.
- [3] Zhihan Chen, Yuhuan Zhao, Yijie Zhu, and Xinyu Yao. When identities collapse: A stress-test benchmark for multi-subject personalization. *arXiv preprint arXiv:2603.26078*, 2026.
- [4] Rinon Gal, Yuval Alaluf, Yuval Atzmon, Or Patashnik, Amit H Bermano, Gal Chechik, and Daniel Cohen-Or. An image is worth one word: Personalizing text-to-image generation using textual inversion. In *International Conference on Learning Representations*, 2023.
- [5] Alain Hore and Djemel Ziou. Image quality metrics: Psnr vs. ssim. In *2010 20th international conference on pattern recognition*, pages 2366–2369. IEEE, 2010.
- [6] Kaiyi Huang, Kaiqiang Sun, Jian Enzweiler, et al. T2i-compbench: A comprehensive benchmark for open-world compositional text-to-image generation. In *Advances in Neural Information Processing Systems*, 2023.
- [7] Yuval Kirstain, Adam Polyak, Uriel Singer, Shahbuland Matiana, Joe Penna, and Omer Levy. Pick-a-pic: An open dataset of user preferences for text-to-image generation. In *Advances in Neural Information Processing Systems*, 2023.
- [8] Divake Kumar, Sina Tayebati, Devashri Naik, Ranganath Krishnan, and Amit Ranjan Trivedi. Vlm judges can rank but cannot score: Task-dependent uncertainty in multimodal evaluation. *arXiv preprint arXiv:2604.25235*, 2026.
- [9] Nupur Kumari, Bingliang Zhang, Richard Zhang, Eli Shechtman, and Jun-Yan Zhu. Multi-concept customization of text-to-image diffusion. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 1931–1941, 2023.
- [10] Zhen Li, Mingdeng Cao, Xintao Wang, Zhongang Qi, Ming-Ming Cheng, and Ying Shan. Photomaker: Customizing realistic human photos via stacked id embedding. *arXiv preprint arXiv:2312.04461*, 2024.
- [11] Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Sypetkowski, Vincent Lempereur, Armand Guzmão, Armand Joulin, and Piotr Bojanowski. Dinov2: Learning robust visual features without supervision. *arXiv preprint arXiv:2304.07193*, 2023.
- [12] Yuta Oshima, Daiki Miyake, Kohsei Matsutani, Yusuke Iwasawa, Masahiro Suzuki, Yutaka Matsuo, and Hiroki Furuta. Multibanana: A challenging benchmark for multi-reference text-to-image generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026. *arXiv preprint arXiv:2511.22989*.
- [13] Yuang Peng, Yuxin Cui, Haomiao Su, Mingzhen Ma, Wenqi Fang, Ting Cheng, Guanzhong Feng, Yu Hu, and Zhen Zhao. Dreambench++: A human-aligned benchmark for personalized image generation. *arXiv preprint arXiv:2406.16855*, 2024.
- [14] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. *International conference on machine learning*, pages 8748–8763, 2021.
- [15] Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman. Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 22500–22510, 2023.

- [16] Wenqing Tian, Hanyi Mao, Zhaocheng Liu, Lihua Zhang, Qiang Liu, Jian Wu, and Liang Wang. Multibind: A benchmark for attribute misbinding in multi-subject generation. *arXiv preprint arXiv:2603.21937*, 2026.
- [17] Shulei Wang, Longhui Wei, Xin He, Jianbo Ouyang, Hui Lu, Zhou Zhao, and Qi Tian. Psr: Scaling multi-subject personalized image generation with pairwise subject-consistency rewards. *arXiv preprint arXiv:2512.01236*, 2025.
- [18] X Wang, Siming Fu, Qihan Huang, Wanggui He, and Hao Jiang. Ms-diffusion: Multi-subject zero-shot image personalization with layout guidance. *arXiv preprint arXiv:2406.07209*, 2024.
- [19] Zhou Wang, Alan C Bovik, Hamid R Sheikh, and Eero P Simoncelli. Image quality assessment: from error visibility to structural similarity. *IEEE transactions on image processing*, 13(4):600–612, 2004.
- [20] Zhou Wang, Eero P Simoncelli, and Alan C Bovik. Multiscale structural similarity for image quality assessment. In *The thirty-seventh asilomar conference on signals, systems & computers, 2003*, volume 2, pages 1398–1402. Ieee, 2003.
- [21] Chenyuan Wu, Pengfei Zheng, Ruiran Yan, Shitao Xiao, Xin Luo, Yueze Wang, Wanli Li, Xiyan Jiang, Yexin Liu, Junjie Zhou, et al. Omnigen2: Exploration to advanced multimodal generation. *arXiv preprint arXiv:2506.18871*, 2025.
- [22] Xiaoshi Wu, Yiming Hao, Keqiang Sun, Yixiong Chen, Feng Zhu, Rui Zhao, and Hongsheng Li. Human preference score v2: A solid benchmark for evaluating human preferences of text-to-image synthesis. *arXiv preprint arXiv:2306.09341*, 2023.
- [23] Guangxuan Xiao, Tianwei Yin, William T Freeman, Fredo Durand, and Song Han. Fastcomposer: Tuning-free multi-subject image generation with localized attention. In *International Journal of Computer Vision*, 2024.
- [24] Jiazheng Xu, Xiao Liu, Yien Wu, Yuxuan Tong, Qinkai Li, Ming Ding, Jie Tang, and Yuxiao Dong. Imagereward: Learning and evaluating human preferences for text-to-image generation. In *Advances in Neural Information Processing Systems*, 2023.
- [25] Tang Xu, Wenbin Wang, Alin Zhong, et al. Customized human object interaction image generation. *arXiv preprint arXiv:2508.19575*, 2025.
- [26] Hu Ye, Jun Zhang, Sibio Liu, Xiao Han, and Wei Yang. Ip-adapter: Text compatible image prompt adapter for text-to-image diffusion models. *arXiv preprint arXiv:2308.06721*, 2023.
- [27] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 11975–11986, 2023.

407 A Appendix: Human Annotation Observations and Failure Mode Analysis

408 During the human annotation process for the gold test set, annotators systematically recorded
 409 recurring failure patterns across evaluated image pairs. These observations surfaced three consistent
 410 and practically significant findings, which we document here to inform future dataset construction
 411 and model development efforts.

412 **Failure Rates Increase with Subject Count.** Annotators consistently found that images depicting
 413 a larger number of subjects were significantly harder to evaluate correctly, and that generative models
 414 performed more poorly on such cases. The most common failure modes in high-subject-count scenes
 415 were subject omission, where one or more individuals were missing from the generated output, and
 416 appearance drift, where the physical attributes of individual subjects degraded as the total number
 417 of subjects increased. This trend was particularly pronounced for visually similar subjects (e.g.,
 418 individuals wearing comparable outfits or sharing similar physical features), where models frequently
 419 confused or merged identities. Facial attributes and clothing details were identified as the two most
 420 error-prone visual channels under these conditions.




Figure 6: **Failure Rates Increase with Subject Count** A representative 6-subject, 2-object prompt (ID: 6905) evaluated across four generative models. All models exhibit Existence failures, with subject omission ranging from 1 missing subject (FLUX2, Seedream-4.5, GPT) to 3 missing subjects (GLM).

Existence Failures Frequently Co-occur with Appearance Failures. A substantial proportion of annotated samples exhibited a systematic co-occurrence between Existence and Appearance failures. When a model fails to render all requested subjects as distinct entities, it often compensates by blending the features of missing subjects into the remaining ones, for instance, combining the facial identity of subject *A* with the wardrobe of subject *B* into a single generated figure. This entanglement means that Existence failures are rarely isolated: they tend to cascade into cross-subject feature bleeding, inflating Appearance failure rates in a correlated rather than independent manner. This observation motivates treating the three diagnostic dimensions as structured but not fully orthogonal signals in practice, despite their conceptual independence.




Figure 7: **Existence Failures Frequently Co-occur with Appearance Failures** A representative 4-subject prompt (ID: 32062) generated by Seedream, the best-performing model on this sample. The missing subject (Man in Black Suit) is not simply absent: its visual attributes, the black tie and formal shirt, are redistributed onto the surviving Bomber Jacket figure, producing a hybrid identity.

Multi-Action Overload Leads to Subject Deformation Annotators consistently observed that when a single subject is assigned more than one concurrent action or interaction, generative models struggle to faithfully execute all specified behaviors. This overload manifests in two characteristic failure patterns. First, *figure splitting*: the model resolves the conflict by duplicating or fragmenting the subject, distributing distinct actions across two visually similar but non-identical figures, effectively cloning the identity to satisfy each action independently. Second, *action dropout*: the model selectively ignores or partially executes one of the assigned actions, producing a subject that satisfies only a subset of the prompt specification. Both failure modes were reproducible across multiple models and prompt configurations, suggesting that multi-action binding exceeds the compositional capacity of current generative architectures rather than reflecting prompt ambiguity.




Figure 8: **Multi-Action Overload Leads to Subject Deformation.** When a subject is assigned more than one concurrent action, models either clone the subject to satisfy each action independently (Figure Splitting, ID: 61186) or silently discard one of the assigned actions (Action Dropout, ID: 60869). Both failure modes are observed in Nano Banana, a better performing model in these cases.

440 B Appendix: Summarized Table of Related Work

Existing benchmarks for personalized and multi-subject image generation have made significant progress, yet they collectively leave a critical evaluation gap that MIB is designed to address. As summarized in Table 2, earlier text-to-image benchmarks such as T2I-CompBench [6] entirely lack reference image support, making them unable to measure subject identity fidelity. Single-subject benchmarks like DreamBench [15] and DreamBench++ [13] introduced reference-based evaluation but restrict testing to one concept at a time, leaving multi-subject binding behavior unexamined. More recent efforts including XVerseBench [2], PSRBench [17], OmniContext [21], MultiBanana [12], and MultiHuman-Testbench [1] extend evaluation to multiple subjects, yet none provides factorized relation-level stress testing that isolates failure modes arising from physical interaction, occlusion, and identity leakage under increasing subject count. MultiBind [16] targets attribute-level misbinding diagnostics but does not consider the geometric and occlusion-driven binding failures that emerge in contact-rich scenes. Crucially, no prior benchmark provides human pairwise gold evaluation labels, which are essential for validating automated metric reliability in the multi-subject setting. MIB addresses all of these gaps simultaneously: it introduces factorized relation-level binding stress tests across up to eight subjects, pairs scalable silver supervision with human gold pairwise preference labels, and provides the first diagnostic framework specifically designed to expose and quantify identity collapse and appearance leakage in complex multi-subject generation.

Benchmark	Year	Ref. Images	Multi-subj.	Stress / Eval. Design	Human Gold Eval.	Pref/Diag. Supervision	Main Gap
T2I-CompBench [6]	2023	×	Text-only	Text-only composition	×	×	Cannot measure reference identity, identity collapse, or appearance leakage.
DreamBench [15]	2023	✓	×	Single-subject prompts	×	×	No multi-subject binding or interaction evaluation.
DreamBench++ [13]	2024	✓	×	Single-subject, human-aligned automated eval.	×	×	Broader single-subject evaluation, but no high-subject-count binding stress.
CustomConcept101 [9]	2023	✓	Limited	Multi-concept composition	×	×	Not designed for contact-rich or occlusion-driven identity binding.
XVerseBench [2]	2025	✓	✓	Single / dual / triple subjects	×	×	Limited high-subject-count and relation-level stress factorization.
PSRBench [17]	2025	✓	✓	Broad capability subsets	×	Reward metrics	Capability-oriented; not a factorized diagnostic benchmark for interaction/occlusion binding failures.
OmniContext [21]	2025	✓	✓	In-context character / object / scene eval.	×	Automated judge scores	Broad in-context consistency benchmark, but lacks human gold and factorized relation-level binding stress.
MultiBanana [12]	2025	✓	✓	Multi-reference challenges	×	VLM-based scores	Covers multi-reference difficulty, but not hierarchical interaction/occlusion diagnosis or human pairwise gold labels.
MultiHuman-Testbench [1]	2025	✓	✓	Multi-human actions / faces	×	Metric suite	Restricted to human subjects; does not cover human-object or object-object binding.
MultiBind [16]	2026	✓	✓	Attribute / slot misbinding	×	Diagnostic metrics	Focuses on attribute-level misbinding rather than factorized physical interaction and occlusion stress.
MIB (Ours)	2026	✓	✓	Factorized relation-level binding stress	✓	✓	–

Table 2: Comparison with existing benchmarks. DreamBench and CustomConcept101 refer to evaluation sets from DreamBooth and Custom Diffusion respectively. MIB uniquely targets factorized relation-level binding stress with both scalable silver supervision and human gold evaluation.

C Appendix: SOP Prompt

You are an expert judge for multi-subject personalized image generation.

You will receive:

1. The original generation prompt.
2. Reference images for each subject mentioned in the prompt.
3. Candidate image A.
4. Candidate image B.

Your job consists of two parts:

Part 1: Independent Evaluation

Evaluate candidate A and candidate B independently. For each image, strictly evaluate:

- Existence: Are ALL the required subjects in the references physically present? (Generic versions count as 1; completely missing counts as 0).
- Appearance: Are the subjects rendered WITHOUT structural deformations? Do the generated subjects strictly match the references across ALL features: Face/Skin tone, Hair style, Top clothing, AND Bottoms (pants)?
- Interaction: Does the generated image reflect the described actions, maintain correct proportions, and follow realistic spatial/physical laws?

Part 2: Overall Comparison


Compare Candidate A and B and decide the overall winner.

OUTPUT RULES:


- Use ONLY the provided prompt and images.
- All dimensional scores ('*_existence', '*_appearance', '*_interaction') must be strictly binary: 1 (fully correct) or 0 (any issue exists).
- 'winner' must be either "A" or "B" (Never "Tie").
- Return valid JSON ONLY. Absolutely NO markdown formatting.
- CRITICAL: To focus on errors, the JSON MUST begin with 'a_flaw_log' and 'b_flaw_log'. DO NOT list perfect subjects. ONLY list subjects with errors and explicitly state what is wrong (Missing, Wrong Face/Hair/Skin, Wrong Top,

490 Wrong Pants, Action/Spatial/Physics Error). If an image is 100% perfect,
 491 write "None".
 492
 493 SCORING NOTES & SPECIFIC EDGE CASES:
 494 - Appearance (Strict Identity & Wardrobe): You MUST check the whole person. If a
 495 subject has the correct top but wrong pants, wrong hair, or mismatched skin
 496 tone/face compared to the reference, score Appearance as 0 and log the
 497 specific flaw (e.g., "Woman red hoodie: Wrong face/skin color",
 498 "Man flannel shirt: Wrong pants color").
 499 - Appearance (Mangled/Melted): Any structural artifacts (extra limbs, mangled
 500 faces) = Appearance 0.
 501 - Existence vs Appearance: If requested "man denim jacket" is drawn as a generic
 502 man in a grey sweater, Existence is 1, Appearance is 0.
 503 - Interaction (Partial): Evaluate interactions based ONLY on rendered subjects.
 504 If an interaction requires a missing subject, Interaction is 0.
 505 - Interaction (Physics, Environment & Scale): Obvious physical violations
 506 (e.g., objects floating in mid-air without support) OR severe scale distortions
 507 between objects (e.g., abnormally giant bread, miniature people) = Interaction 0.
 508 HOWEVER, the presence of unprompted but logical supporting structures
 509 (e.g., tables, pedestals) used to hold objects is PERMITTED and should NOT be
 510 penalized.
 511
 512 EXPECTED JSON SCHEMA:
 513 {
 514 "a_flaw_log": "Middle eastern man: Appearance (wears green sweater instead of
 515 beige shirt). Man flannel shirt: Appearance (wrong pants color). Woman red
 516 hoodie: Appearance (mismatched face/skin/hair). Interaction: Man is not
 517 occluded by the woman; The burger is floating in mid-air.",
 518 "b_flaw_log": "Middle eastern man: Missing. Man bomber jacket: Missing. DSLR
 519 camera: Missing. Interaction: Man is not walking towards/staring at the woman;
 520 Fails due to missing interaction targets.",
 521 "a_existence": 1,
 522 "a_appearance": 0,
 523 "a_interaction": 0,
 524 "b_existence": 0,
 525 "b_appearance": 0,
 526 "b_interaction": 0,
 527 "winner": "A"
 528 }


529 D Appendix: Demos of Generated Training Labels




(a) Action and Target Errors




(b) Physics and Spatial Violations



(c) Appearance and Occlusion Errors



(d) Scale Distortion and Color Mismatch



(e) Robust Diagnostics in Highly Complex (6-Subject) Scenes

Figure 9: Qualitative examples of strict cross-model consensus. The externalized flaw logs demonstrate that both VLMs effectively utilize the SOP to catch fine-grained errors in appearance, physical realism, interaction semantics, and complex multi-subject scaling.

Table 3: Training Hyperparameters (configuration shown for the 0.8B variant; 2B/4B variants use the same setup with adjusted batch size to fit memory constraints).

Parameter	Value
Base Model	Qwen3.5-0.8B (via unsloth)
Fine-tuning Strategy	Layer-only unfreezing (top 4 layers)
Optimizer	AdamW
Learning Rate	2×10^{-5}
Weight Decay	0.01
LR Scheduler	Cosine schedule with warmup
Warmup Ratio	0.03
Batch Size (per device)	4
Gradient Accumulation Steps	8
Target Effective Batch Size	16
Max Epochs	3 (dynamic auto-scaling: 2–6)
Gradient Clipping Norm	1.0
Image Resolution	512×512
Global Random Seed	3407

Table 4: Multi-task Loss Configuration

Component	Setting
Preference Ranking Loss	Margin Ranking Loss (margin = 0.1)
Ranking Loss Weight (α)	1.0
Diagnostic Classification Loss	BCEWithLogitsLoss
Classification Loss Weight (β)	0.5
Early Stopping	Patience: 2 epochs, min_delta = 10^{-3}

Table 5: Inference and Evaluation Setup

Setting	Detail
MIE Inference Precision	Default checkpoint precision (bf16/fp16)
Evaluation Resolution	512×512 (preserving aspect ratio)
Metric 1: Preference Accuracy	Binary choice based on scoring head logits
Metric 2: Diagnostic Output	Weakest dimension via Sigmoid from classification head

Table 6: Baseline Models Configuration

Model	Configuration
CLIP Baseline	openai/clip-vit-base-patch32 Text Truncation: Enabled Scoring: Cosine similarity between text embedding and pooled image embedding
DINO Baseline	facebook/dinov2-base Scoring: Cosine similarity between the average normalized subject reference embeddings and generated image embeddings

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The contributions and scope are clearly stated and discussed in abstract and introduction.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The limitations are discussed in Section 5.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: This field does not apply to our study.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The code and data are provided in the supplementary file, thus, the results are reproducible.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The benchmark data and code access links are available in the supplementary files. MIB Benchmark and MIBE Repo also available here.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: In Section 3, we discuss the construction of training and testing sets. Section 4 and Appendix shows the details of evaluator model architecture and hyperparameters.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [N/A]

Justification: This field does not apply to our study.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).

- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: In Section 4 and appendix, we show the configuration of the computation resources.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The research conforms with the NeurIPS Code of Ethics. No personally identifiable information was collected, and all human annotation was conducted under an anonymous pairwise protocol.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Shown in Section 1 and 6, MIBE provides the community with a 60K silver training corpus and a 4K human-labeled gold benchmark, enabling more rigorous evaluation of multi-subject personalized image generation. MIE offers a significantly more reliable metric for complex multi-reference scenarios where existing evaluators collapse, supporting more accurate model comparison and alignment research.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.

- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: MIB and MIE are evaluation resources rather than generative models, and do not directly enable the creation of harmful content. No additional safeguards beyond standard CC BY 4.0 licensing terms are deemed necessary.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: In Section 3 and 4, we cited and disclosed that we used LLMs to generate prompts, used VLMs to generate the MIB-silver set labels, and we leveraged backbone VLMs to fine tune a multi-subject image evaluator.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.

- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [\[Yes\]](#)

Justification: In the Supplementary materials, we show that the benchmark dataset and code are available in the provided MIB Benchmark and MIBE Repo.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [\[N/A\]](#)

Justification: All annotators in this study are the authors of this paper. Therefore, this work does not involve external crowdsourcing or human subject research.

Guidelines:

- The answer [\[N/A\]](#) means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [\[N/A\]](#)

Justification: Our paper does not involve crowdsourcing nor research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: In Section 3 and 4, we disclose that we used LLMs to generate prompts, used VLMs to generate the MIB-silver set labels, and we leveraged backbone VLMs to fine tune a multi-subject image evaluator.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.