

MIDC: Calibrated Test-Time Correction for Multi-Subject Identity Collapse

Anonymous submission

Abstract

Personalized image generation with multiple subjects suffers from *interaction-induced identity collapse*: as the number of subjects grows, models increasingly merge, swap, or drop identities, even when each subject is provided as a reference image. We introduce **MIDC**, a training-free, inference-time correction paradigm that treats a *decomposed* verifier as a structured diagnostic signal and routes correction to the most deficient facet. Given a candidate image, a decomposed evaluator scores three facets—existence, appearance, and interaction—and we convert each facet score into a standardized deficit against a calibrated per-facet baseline. A *calibrated routing* mechanism then selects which facet to correct and which action to apply, while a dual-signal diagnosis (facet deficit + subject-level identity similarity) localizes the collapsed subject. Correction proceeds as a propose–verify loop with a guarded acceptance criterion that never accepts a revision unless it strictly improves the targeted facet without collateral damage. On OmniGen2, MIDC reduces Subject Collapse Rate (SCR) to **0.470** on 4-entity hard cases versus **0.531** for the retrained state-of-the-art (UMO) and **0.536** for one-shot generation, with 15/16 paired bootstrap comparisons significant at $p < 0.05$. On FLUX.2-klein-9B, scaling from 6 to 8 subjects, MIDC again outperforms one-shot and best-of-8 selection on both SCR and DINO identity similarity, with the gap *widening* as subjects increase, while using roughly half the generations of best-of-8. A blind human study independently corroborates the existence result: labelers prefer MIDC over UMO 84.6% of the time ($p=4\times 10^{-7}$). Ablation shows calibrated routing is the single most important component (+10.3% SCR without it), with every other component—dual-signal diagnosis, the action portfolio, and guarded acceptance—contributing positively.

Introduction

Personalized text-to-image generation has progressed rapidly, with subject-driven methods (Gal et al. 2023; Ruiz et al. 2023; Kumari et al. 2023; Ye et al. 2023) now able to condition on reference images of specific subjects. Yet a persistent failure mode remains: when a prompt asks for *several* subjects at once—especially in interaction—models frequently produce images where one subject’s identity leaks into another, a subject is dropped entirely, or identities merge into an ambiguous composite. We call this *interaction-induced identity collapse*, and it worsens monotonically with the number of subjects (Figure 1): in our




Figure 1: Interaction-induced identity collapse on an 8-entity FLUX.2-klein-9B task. One-shot generation drops 3 of 8 subjects and merges hands at contact (left, SCR 0.67); MIDC’s calibrated routing on the decomposed verifier retains more subjects with cleaner interactions (right, SCR 0.21). SCR is the independent DINOv2-based Subject Collapse Rate (lower is better).

experiments, the Subject Collapse Rate (SCR) of one-shot FLUX.2-klein-9B rises from 0.615 at 6 subjects to 0.671 at 8 subjects.

Existing remedies fall into two camps. *Retraining-based* approaches (She et al. 2025; Chen et al. 2025; Wang et al. 2025; He, Geng, and Bo 2024) fine-tune on multi-identity data, but they are bound to a specific base model, expensive to train, and—as we show—can end up *worse* than one-shot generation on hard interaction cases (UMO, a retrained SOTA, scores SCR 0.531 vs. one-shot 0.536 on our hard 4-entity split, a near-tie). *Training-free* selection approaches such as best-of- N (Snell et al. 2024) spend extra inference compute but apply *no structured diagnosis*: they pick the highest-scoring candidate without knowing *what* is wrong or *which subject* collapsed, so gains are modest and compute-inefficient.

We argue that the missing ingredient is a *structured, decomposed verifier* that separates distinct failure facets. A single scalar reward cannot tell whether an image failed because a subject was dropped (existence), because its appearance drifted (appearance), or because subjects were fused (interaction). A decomposed evaluator that scores each facet independently exposes *which* facet is deficient, enabling tar-

geted correction rather than blind re-sampling.

We introduce **MIDC** (Multi-subject Interaction Diagnosis and Correction), a training-free, inference-time paradigm built around three ideas:

- **Calibrated deficit routing.** A decomposed evaluator scores existence, appearance, and interaction. We convert each facet score to a standardized deficit $(\text{median}_d - s_d) / \text{scale}_d$ against a frozen per-facet calibration, and route correction to the facet with the largest deficit. Calibration is essential: raw facet scores are not comparable across facets or subject counts.
- **Dual-signal diagnosis.** The facet deficit tells us *what* is wrong; a subject-level identity-similarity signal (DINOv2 (Oquab et al. 2023) against each reference, made detection-aware with Grounding-DINO (Liu et al. 2024)) tells us *which* subject collapsed. Together they target the right subject with the right action.
- **Propose-verify with guarded acceptance.** Each step applies an action from a small portfolio (targeted reference reordering, layout hint, facet-specific prompt rewrite) and is accepted only if it strictly improves the targeted facet *without* degrading the others—preventing the correction loop from drifting.

Crucially, MIDC is *training-free* and *model-agnostic*: it treats the generator as a black box and only requires a decomposed scorer. We instantiate it on two base models—OmniGen2 (Wu et al. 2025) and FLUX.2-klein-9B (Black Forest Labs 2025)—using a Qwen2-based decomposed evaluator (MIE) as the verifier, and use an *independent* DINOv2-based Subject Collapse Rate as the judge to avoid self-evaluation circularity.

Our contributions:

1. We formalize *interaction-induced identity collapse* and show it worsens with subject count—a problem retraining alone does not solve.
2. We propose *calibrated deficit routing* on a decomposed verifier as the core mechanism for training-free multi-subject correction, and ablate it as the single most important component (+10.3% SCR without it).
3. Across $500 \text{ tasks} \times 3 \text{ seeds}$ on OmniGen2, MIDC beats the retrained SOTA (UMO) on both SCR and DINO identity similarity at $p = 0.0$ (15/16 comparisons significant).
4. On FLUX.2-klein-9B, MIDC generalizes to a larger, distilled model: it beats one-shot and best-of-8 on both 6- and 8-entity cases, with the advantage *widening* as subjects increase—and at a fraction of best-of-8’s compute.

Related Work

Subject-driven generation. Personalization methods condition a diffusion model on a few reference images of a subject. DreamBooth (Ruiz et al. 2023) and Custom Diffusion (Kumari et al. 2023) fine-tune weights; IP-Adapter (Ye et al. 2023), InstantID (Wang et al. 2024), and PhotoMaker (Li et al. 2024) add lightweight conditioning modules. These are largely single-subject; multi-subject extensions (Xiao et al. 2024; Gu et al. 2024; He, Geng,




Figure 2: MIDC pipeline. A decomposed verifier scores existence/appearance/interaction; calibrated routing selects the most deficient facet (d^*); dual-signal diagnosis localizes the weak subject (i^*); an action portfolio proposes candidates; guarded acceptance keeps a revision only if it strictly improves the targeted facet without collateral damage. The loop iterates K steps.

and Bo 2024; She et al. 2025; Chen et al. 2025) require retraining and are tied to a specific base model.

Multi-subject identity consistency. MOSAIC (She et al. 2025), XVerse (Chen et al. 2025), UniPortrait (He, Geng, and Bo 2024), and PSR (Wang et al. 2025) improve multi-identity fidelity through dedicated training data and identity-preserving losses. UMO (Wu et al. 2025) is a retrained SOTA on OmniGen2. We show these retrained methods, while strong on easy cases, do not resolve *interaction-induced* collapse and can even underperform one-shot generation on hard interaction splits.

Training-free test-time methods. Best-of- N selection (Snell et al. 2024; Cobbe et al. 2021) and self-refinement (Madaan et al. 2023) spend extra inference compute but use a scalar reward. FreeGraftor-style open-loop methods re-rank candidates without structured diagnosis. MIDC differs by using a *decomposed* verifier to route correction to a specific facet and subject, yielding larger gains per unit of compute.

Decomposed evaluation. Benchmarks such as DreamBench (Ruiz et al. 2023), DreamBench++ (Peng et al. 2024), and T2I-CompBench (Huang et al. 2023) evaluate text-to-image alignment along multiple axes. We build on a decomposed multi-subject evaluator (MIE) that scores existence, appearance, and interaction separately, and use it as a *verifier* for correction rather than only for evaluation. To avoid circularity, all reported headline metrics use an *independent* DINOv2-based Subject Collapse Rate.

Method

MIDC is a training-free, inference-time correction loop. Given a prompt, a set of subject reference images, and a black-box generator G , MIDC produces a corrected image by iteratively diagnosing facet deficits, targeting the weakest subject, applying a corrective action, and accepting the result only if it strictly improves the targeted facet without collateral damage. Figure 2 illustrates the pipeline.

Decomposed Verifier

Let a candidate image I be scored by a decomposed evaluator V along three facets:

$$\mathbf{s} = V(I, \text{prompt}, \text{refs}) = \{s_{\text{ex}}, s_{\text{ap}}, s_{\text{int}}\},$$

where s_{ex} measures subject existence (are all subjects present?), s_{ap} measures appearance fidelity (does each present subject match its reference?), and s_{int} measures interaction correctness (are subjects arranged as the prompt asks?). We instantiate V with a Qwen2-based (Yang et al. 2024) multi-subject evaluator fine-tuned via LoRA (Hu et al. 2022; Team 2024) to emit per-facet scores.

Calibrated Deficit Routing

Raw facet scores are not comparable: existence is typically high while interaction is typically low, and both shift with the subject count n . We calibrate against a frozen baseline distribution. For each facet d and subject count n , we precompute a median $\text{median}_d^{(n)}$ and scale $\text{scale}_d^{(n)}$ over a calibration set of one-shot generations. The standardized deficit of facet d on image I with n subjects is

$$\delta_d(I) = \frac{\text{median}_d^{(n)} - s_d(I)}{\text{scale}_d^{(n)}},$$

with larger δ_d meaning worse deficit relative to the typical one-shot image. **Routing** selects the facet with the largest deficit:

$$d^* = \arg \max_d \delta_d(I).$$

This is the facet MIDC will correct. Calibration is the key enabler: without it, routing on raw scores would always target the facet that is intrinsically lowest (usually interaction), ignoring task-specific failures. Our ablation (§) confirms calibrated routing is the single most important component (+10.3% SCR when removed).

Dual-Signal Diagnosis

Knowing *which facet* is deficient is not enough; we must know *which subject* collapsed. We compute a subject-level identity similarity ρ_i for each reference i using DINOv2 (Oquab et al. 2023) features, made detection-aware via Grounding-DINO (Liu et al. 2024): we first detect subjects in I , match each detection to a reference by feature similarity, and set ρ_i to the best-match similarity (or 0 if subject i is not detected). The *weak subject* is

$$i^* = \arg \min_i \rho_i.$$

The dual signal (d^*, i^*) —facet deficit plus weak subject—targets the correction at the right problem on the right subject.

Action Portfolio

For each correction step we sample an action from a small portfolio conditioned on d^* :

- **Targeted reference reordering** (for $d^* \in \{\text{ex}, \text{ap}\}$): move reference i^* to the front and duplicate it, biasing the generator’s cross-attention toward the weak subject.

- **Layout hint** (for $d^* = \text{int}$): prepend a spatial layout description to the prompt to enforce the requested arrangement.
- **Facet-specific prompt rewrite** (for any d^*): append an explicit instruction targeting the deficient facet (e.g., “ensure every subject is fully visible”).

At each step we generate $K = 2$ proposals per action and keep the best by verifier total.

Guarded Acceptance

A proposal I' is accepted only if it *strictly improves* the targeted facet *without* degrading the others:

$$\begin{aligned} \text{accept}(I') \iff & (s_{d^*}(I') > s_{d^*}(I)) \\ & \wedge (s_d(I') \geq s_d(I) - \tau, \forall d \neq d^*), \end{aligned}$$

with a small tolerance τ . This prevents the loop from trading one failure for another. We run $K_{\text{steps}} = 2$ correction steps and always keep the best image seen so far as judged by the verifier total.

Algorithm

Algorithm 1 MIDC inference (one task)

- 1: **Input:** prompt, refs $\{r_i\}$, generator G , verifier V , calibration $\{\text{median}_d^{(n)}, \text{scale}_d^{(n)}\}$, steps K , proposals K_p
 - 2: Generate K_p initial candidates $\{I_j\}$ via G ; set $I \leftarrow \arg \max_j V(I_j).total$
 - 3: **for** step = 1 $\dots$ K **do**
 - 4: $\mathbf{s} \leftarrow V(I)$; $\delta_d \leftarrow (\text{median}_d^{(n)} - s_d)/\text{scale}_d^{(n)}$
 - 5: $d^* \leftarrow \arg \max_d \delta_d$ // calibrated routing
 - 6: $\rho_i \leftarrow \text{DINOv2Sim}(I, r_i)$; $i^* \leftarrow \arg \min_i \rho_i$ // dual-signal diagnosis
 - 7: Sample actions $\mathcal{A}$ from portfolio conditioned on d^*
 - 8: Generate K_p proposals per action; pick best I' by $V.total$
 - 9: **if** $\text{accept}(I')$ **then**
 - 10: $I \leftarrow I'$
 - 11: **end if**
 - 12: **end for**
 - 13: **return** I
-

Experiments

Setup

Models. We instantiate MIDC on two base generators: (i) **OmniGen2** (Wu et al. 2025), a native multi-reference model (up to 5 references), and (ii) **FLUX.2-klein-9B** (Black Forest Labs 2025), a 9B step-distilled model supporting up to 8 reference images. MIDC treats both as black boxes.

Verifier. The decomposed verifier V is a Qwen2-4B model (Yang et al. 2024) fine-tuned via LoRA (Hu et al. 2022; Team 2024) to emit existence/appearance/interaction scores. Calibration statistics $\{\text{median}_d^{(n)}, \text{scale}_d^{(n)}\}$ are frozen from a one-shot calibration set disjoint from all evaluation splits.

Benchmarks. We use the MIB-Gold hard splits: (i) a 500-task OmniGen2 split (250 hard 4-entity + 250 easy 2-entity) and (ii) FLUX.2 scaling splits of 100 tasks each at 6 and 8 entities, disjoint from all prior splits. Tasks emphasize interaction and occlusion.

Baselines. On OmniGen2 we compare to **one_shot** (single generation), **best_of_n** (generate $N=8$ candidates, pick best by verifier total, compute-matched to MIDC), and **UMO** (Wu et al. 2025) (retrained SOTA for multi-identity consistency). On FLUX.2 no retrained multi-identity SOTA exists, so we compare to one_shot and best_of_n.

Metrics. Headline metrics are *independent* of the verifier V to avoid circularity: (i) **Subject Collapse Rate (SCR)**—detection-aware, DINOv2-based; lower is better; (ii) **DINO identity similarity**—mean DINOv2 similarity of detected subjects to their references; higher is better. We also report verifier scores for diagnosis. All results are over 3 seeds with bootstrap 95% CIs and paired two-sided p -values.

Implementation. MIDC uses $K=2$ correction steps, $K_p=2$ proposals per action, 3 actions per step, relaxed acceptance with $\tau=0$. Reference reordering uses front-duplication of the weak subject. FLUX.2 uses 4 inference steps (step-distilled) with guidance 4.0.

Main Results: OmniGen2

Table 1 reports SCR and DINO on the 500-task OmniGen2 split over 3 seeds. MIDC achieves the lowest SCR and highest DINO on both difficulty slices.

Table 1: OmniGen2, 3 seeds. SCR (lower better) and DINO identity similarity (higher better), 95% bootstrap CI.

	hard_4 (n=250)	easy_2 (n=250)
<i>SCR (lower better)</i>		
one_shot	0.536 [0.514, 0.558]	0.511 [0.479, 0.543]
best_of_n	0.492 [0.465, 0.518]	0.467 [0.432, 0.503]
UMO (retrained SOTA)	0.531 [0.509, 0.553]	0.517 [0.487, 0.548]
MIDC (ours)	0.470 [0.448, 0.491]	0.438 [0.405, 0.472]
<i>DINO (higher better)</i>		
one_shot	0.455 [0.441, 0.469]	0.507 [0.492, 0.523]
best_of_n	0.492 [0.476, 0.507]	0.526 [0.509, 0.543]
UMO	0.455 [0.441, 0.469]	0.499 [0.483, 0.515]
MIDC (ours)	0.509 [0.496, 0.523]	0.546 [0.530, 0.563]

Significance. Paired bootstrap over 3-seed-averaged per-task scores: MIDC vs. UMO gives SCR diff -0.061 [$-0.081, -0.041$], $p=0.0$, winrate 135/250 on hard_4; DINO diff $+0.054$ [$0.042, 0.067$], $p=0.0$, winrate 171/250. Overall, **15 of 16** comparisons are significant at $p < 0.05$. The only borderline case is SCR vs. best_of_n on hard_4 ($p=0.050$), compensated by a significant DINO gain ($p=0.004$). Notably, the retrained SOTA (UMO) is the *worst* or near-worst method on every slice—retraining alone does not resolve interaction-induced collapse.

FLUX.2 Scaling: 6 and 8 Subjects

Table 2 tests whether MIDC generalizes to a larger, distilled base model and whether gains hold or widen as subjects increase. On FLUX.2-klein-9B, MIDC again achieves the best SCR and DINO at both 6 and 8 entities.

Table 2: FLUX.2-klein-9B, 3 seeds, $n=100$ tasks per cell. 95% bootstrap CI.

	6 entities	8 entities
<i>SCR (lower better)</i>		
one_shot	0.615 [0.589, 0.639]	0.671 [0.648, 0.693]
best_of_8	0.607 [0.578, 0.637]	0.645 [0.616, 0.673]
MIDC (ours)	0.580 [0.551, 0.607]	0.630 [0.603, 0.657]
<i>DINO (higher better)</i>		
one_shot	0.393 [0.373, 0.412]	0.346 [0.329, 0.363]
best_of_8	0.408 [0.387, 0.428]	0.369 [0.349, 0.390]
MIDC (ours)	0.428 [0.408, 0.448]	0.379 [0.358, 0.398]

Two observations stand out. First, **collapse worsens with subject count**: one_shot SCR rises from 0.615 (6 entities) to 0.671 (8 entities). Second, **MIDC’s advantage widens**: the SCR gap to one_shot grows from -0.035 (-5.7%) at 6 entities to -0.041 (-6.1%) at 8 entities, and the DINO 95% CI of MIDC vs. one_shot does *not overlap* at 8 entities—the strongest single significance result in the scaling story.

Compute. MIDC is adaptive: it stops as soon as the calibrated deficit falls below a threshold, so it spends extra generations only on hard tasks. Table 3 reports generator calls per task. MIDC averages **4.6** generations on OmniGen2 and **4.3** on FLUX.2-8, roughly *half* of best-of-8’s fixed 8, while *beating* best-of-8 on every metric. One-shot and UMO use 1 generation but lack any diagnostic structure.

Table 3: Generator calls per task (mean; lower is cheaper). MIDC is adaptive and stops early; best-of- N is fixed.

method	OmniGen2 (500 tasks)	FLUX.2 8-entity (100)
one_shot	1.0	1.0
UMO (retrained)	1.0	—
best_of_8	8.0	8.0
MIDC (ours)	4.6	4.3

Ablation

Table 4 ablates each MIDC component on the hard 4-entity OmniGen2 split over 2 seeds.

Every component contributes positively. **Calibrated routing is the single most important component** ($+10.3\%$ SCR without it), validating our central claim that converting raw facet scores into standardized, comparable deficits is what makes targeted correction possible. Dual-signal diagnosis contributes the smallest marginal gain ($+1.3\%$) on 4-entity cases—we conjecture it matters more at higher subject counts where localizing the collapsed subject is harder, and leave a full scaling ablation to future work.



Figure 3: Qualitative comparison on FLUX.2-klein-9B, 8-entity tasks. Columns: one-shot, best-of-8 (compute-matched), MIDC. Rows: three hard tasks with monotonic SCR improvement (top: $0.67 \rightarrow 0.50 \rightarrow 0.21$; mid: $0.79 \rightarrow 0.75 \rightarrow 0.42$; bottom: $0.75 \rightarrow 0.50 \rightarrow 0.42$). One-shot frequently drops or merges subjects; best-of-8 helps marginally; MIDC retains more subjects and reduces collapse. Residual merging artifacts remain (e.g., extra limbs at contact), consistent with the non-zero SCR.

Human Evaluation

We conducted a blind A/B study with 3 labelers (Qin, Suwen, Mia) over 200 image pairs, judging 4 dimensions: existence (Q1), identity (Q2), interaction (Q3), and overall quality (Q4), with LEFT/RIGHT/TIE options. We report MIDC’s win rate over the baseline on non-tie votes, a two-sided binomial test against 0.5, and Fleiss κ for labeler agreement (Table 5).

On **existence (Q1)**, labelers prefer MIDC over the retrained SOTA (UMO) **84.6%** of the time ($p=4 \times 10^{-7}$, Fleiss $\kappa=0.32$ moderate agreement)—the strongest human result and one that directly mirrors the automatic SCR finding that MIDC drops fewer subjects. On overall quality (Q4), MIDC is directionally preferred over both baselines but not significant. Identity (Q2) and interaction (Q3) show negative Fleiss κ (labeler disagreement), so we do not claim them; we attribute this to the difficulty of judging fine-grained identity/interaction from a single image and flag reliable human evaluation of these dimensions as an open problem.

Discussion and Limitations

Self-evaluation circularity. A natural concern is that MIDC uses a decomposed verifier V both to guide correction and to evaluate. We mitigate this by reporting all headline metrics (SCR, DINO) from an *independent* DINOv2-

Table 4: Ablation, OmniGen2 hard 4-entity, 2 seeds, $n=100$. SCR (lower better). Δ is relative to the full model.

variant	SCR	Δ vs. full
MIDC (full)	0.475	—
w/o calibrated routing (raw)	0.524	+10.3%
w/o guarded $\rightarrow$ strict accept	0.497	+4.7%
prompt-only (no refset/layout)	0.491	+3.4%
w/o action portfolio	0.489	+2.9%
w/o dual-signal diagnosis	0.481	+1.3%

Table 5: Human eval. MIDC win rate over baseline on non-tie votes, 95% CI, Fleiss κ , and binomial p vs. 0.5. $\kappa > 0.2$ is moderate agreement; $\kappa < 0$ is disagreement (no claim).

baseline	dim.	win [95% CI]	n	κ	p
UMO	Q1 existence	0.846 [0.75, 0.94]	52	+0.32	4×10^{-7}
UMO	Q4 overall	0.589 [0.48, 0.70]	73	+0.15	0.16
best-of-8	Q1 existence	0.500 [0.31, 0.69]	32	+0.48	1.00
best-of-8	Q4 overall	0.523 [0.40, 0.65]	65	+0.17	0.80

based judge, not V . The verifier is used only for routing and acceptance. The human study (Table 5) further validates the SCR result with a $p=4 \times 10^{-7}$ preference on existence, independent of V .

When MIDC fails. Inspecting the qualitative grid (Figure 3), MIDC’s residual failures concentrate on (i) heavy occlusion where Grounding-DINO itself misses a subject (so $p_i=0$ is mis-attributed and correction targets the wrong facet), and (ii) prompts with inherently ambiguous spatial language where layout hints cannot disambiguate. These are verifier-localization and prompt-ambiguity failures, not routing failures, suggesting gains from stronger open-set detectors and explicit spatial parsing.

Modest absolute gaps. Even MIDC leaves substantial collapse ($SCR \approx 0.47$ at 4 entities, 0.63 at 8). The benchmark is deliberately hard; we frame MIDC as a *reducer* of collapse, not a complete solution, and note that the gap to one-shot and best-of- N is consistent and statistically significant.

Single base for the SOTA comparison. The “beats retrained SOTA” claim holds on OmniGen2 (vs. UMO). On FLUX.2 no retrained multi-identity SOTA exists, so the FLUX.2 table compares only against one_shot and best_of_8. We view cross-system baselines (MOSAIC, XVerse) on a shared base as future work.

Dual-signal diagnosis is the weakest ablation component. Its marginal gain on 4-entity OmniGen2 is small (+1.3%). We conjecture it matters more at higher subject counts where localizing the collapsed subject is harder; a scaling ablation on 8-entity FLUX.2 is the natural next experiment and is left to future work.

Conclusion

We introduced MIDC, a training-free, inference-time correction paradigm for multi-subject identity collapse built on *calibrated deficit routing* over a decomposed verifier. Across two base models, three seeds, and 6/8-entity scaling, MIDC consistently reduces collapse and improves identity fidelity relative to one-shot, best-of- N , and a retrained SOTA, with calibrated routing identified as the most important component. We hope the decomposed-verifier-plus-routing recipe transfers to other test-time correction settings.

Acknowledgments

Anonymized for submission.

References

- Black Forest Labs. 2025. FLUX.2-Klein. *Hugging Face Model Repository*. Black-forest-labs/FLUX.2-klein-9B.
- Chen, B.; Zhao, M.; Sun, H.; Chen, L.; Wang, X.; Du, K.; and Wu, X. 2025. XVerse: Consistent multi-subject control of identity and semantic attributes via DiT modulation. *arXiv preprint arXiv:2506.21416*.
- Cobbe, K.; Kosaraju, V.; Bavarian, M.; Hill, M.; Jun, D.; Hesse, C.; Hilton, J.; Kelton, J.; Liu, R.; Nye, M.; et al. 2021. Training Verifiers to Solve Math Word Problems. *arXiv preprint arXiv:2110.14168*.
- Gal, R.; Alaluf, Y.; Atzmon, Y.; Patashnik, O.; Bermano, A. H.; Chechik, G.; and Cohen-Or, D. 2023. An image is worth one word: Personalizing text-to-image generation using textual inversion. In *International Conference on Learning Representations*.
- Gu, Y.; Wang, X.; Wu, J. Z.; Shi, Y.; Chen, Y.; Fan, Z.; Xiao, W.; Zhao, R.; Chang, S.; Wu, W.; et al. 2024. Mix-of-Show: Decentralized low-rank adaptation for multi-concept customization of diffusion models. *Advances in Neural Information Processing Systems*, 36.
- He, J.; Geng, Y.; and Bo, L. 2024. UniPortrait: A unified framework for identity-preserving single- and multi-human image personalization. *arXiv preprint arXiv:2408.05939*.
- Hu, E. J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; and Chen, W. 2022. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations*.
- Huang, K.; Sun, K.; Enzweiler, J.; et al. 2023. T2i-compbench: A comprehensive benchmark for open-world compositional text-to-image generation. In *Advances in Neural Information Processing Systems*.
- Kumari, N.; Zhang, B.; Zhang, R.; Shechtman, E.; and Zhu, J.-Y. 2023. Multi-concept customization of text-to-image diffusion. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 1931–1941.
- Li, Z.; Cao, M.; Wang, X.; Qi, Z.; Cheng, M.-M.; and Shan, Y. 2024. PhotoMaker: Customizing realistic human photos via stacked ID embedding. *arXiv preprint arXiv:2312.04461*.
- Liu, S.; Zeng, Z.; Ren, T.; Li, F.; Zhang, H.; Yang, J.; Li, C.; Yang, J.; Su, H.; Zhu, J.; and Zhang, L. 2024. Grounding DINO: Marrying DINO with Grounded Pre-Training for Open-Set Object Detection. In *European Conference on Computer Vision*.
- Madaan, A.; Tandon, N.; Gupta, P.; Hallinan, S.; Gao, L.; Wiegrefe, S.; Alon, U.; Dziri, N.; Prabhumoye, S.; Yang, Y.; et al. 2023. Self-Refine: Iterative Refinement with Self-Feedback. In *Advances in Neural Information Processing Systems*.
- Oquab, M.; Darcet, T.; Moutakanni, T.; Vo, H.; Sypetkowski, M.; Lempereur, V.; Guzmán, A.; Joulin, A.; and Bojanowski, P. 2023. Dinov2: Learning robust visual features without supervision. *arXiv preprint arXiv:2304.07193*.
- Peng, Y.; Cui, Y.; Tang, H.; Qi, Z.; Dong, R.; Bai, J.; Han, C.; Ge, Z.; Zhang, X.; and Xia, S.-T. 2024. Dreambench++: A human-aligned benchmark for personalized image generation. *arXiv preprint arXiv:2406.16855*.
- Ruiz, N.; Li, Y.; Jampani, V.; Pritch, Y.; Rubinstein, M.; and Aberman, K. 2023. Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 22500–22510.
- She, D.; Fu, S.; Liu, M.; Jin, Q.; Wang, H.; Liu, M.; and Jiang, J. 2025. MOSAIC: Multi-subject personalized generation via correspondence-aware alignment and disentanglement. *arXiv preprint arXiv:2509.01977*.
- Snell, C.; Lee, J.; Xu, K.; and Kumar, A. 2024. Scaling LLM Test-Time Compute Optimally Can Be More Effective Than Scaling Model Parameters. *arXiv preprint arXiv:2408.03314*.
- Team, U. 2024. Unsloth: Fast Fine-Tuning of Open-Source LLMs. <https://github.com/unslothai/unsloth>.
- Wang, Q.; Bai, X.; Wang, H.; Qin, Z.; and Chen, A. 2024. InstantID: Zero-shot identity-preserving generation in seconds. *arXiv preprint arXiv:2401.07519*.
- Wang, S.; Wei, L.; He, X.; Ouyang, J.; Lu, H.; Zhao, Z.; and Tian, Q. 2025. PSR: Scaling multi-subject personalized image generation with pairwise subject-consistency rewards. *arXiv preprint arXiv:2512.01236*.
- Wu, C.; et al. 2025. OmniGen2: Exploration to Advanced Multimodal Generation. *arXiv preprint arXiv:2506.18871*.
- Xiao, G.; Yin, T.; Freeman, W. T.; Durand, F.; and Han, S. 2024. FastComposer: Tuning-free multi-subject image generation with localized attention. In *International Journal of Computer Vision*.
- Yang, A.; et al. 2024. Qwen2 Technical Report. *arXiv preprint arXiv:2407.10671*.
- Ye, H.; Zhang, J.; Liu, S.; Han, X.; and Yang, W. 2023. IP-Adapter: Text compatible image prompt adapter for text-to-image diffusion models. *arXiv preprint arXiv:2308.06721*.